

Homework 4 Writeup

Instructions

- Describe any interesting decisions you made to write your algorithm.
- Show and discuss the results of your algorithm.
- Feel free to include code snippets, images, and equations.

Harris Corner Detector

```

1  def get_interest_points(image, descriptor_window_image_width):
2
3      gradient_filter_x = np.array([[ -1, 0, 1]])
4      gradient_filter_y = np.array([[ -1, 0, 1]])
5
6      # Compute gradients by convolving gradient filters with the
       image
7      Ix = cv2.filter2D(src=img_gray, kernel=gradient_filter_x,
       ddepth=-1)
8      Iy = cv2.filter2D(src=img_gray, kernel=gradient_filter_y,
       ddepth=-1)
9
10     # Compute products of derivatives
11     sigma = 1
12     alpha = 0.05
13     threshold = 0.01
14     Ixx = cv2.GaussianBlur(Ix * Ix, (5, 5), sigma)
15     Iyy = cv2.GaussianBlur(Iy * Iy, (5, 5), sigma)
16     Ixy = cv2.GaussianBlur(Ix * Iy, (5, 5), sigma)
17     cornerness = (Ixx*Iyy) - (Ixy**2) - alpha*((Ixx+Iyy)**2)
18     corners = cornerness > threshold * cornerness.max()
19     distance = 1
20
21     #Non Max Suppression
22     for i in range(distance, corners.shape[0] - distance):
23         for j in range(distance, corners.shape[1] - distance):
24             if corners[i, j]:
25                 local_max = cornerness[i - distance:i + distance +
26 1, j - distance:j + distance + 1].max()
27                 if cornerness[i, j] != local_max:
28                     corners[i, j] = False
29
30     # Getting the coordinates of corners
31     y, x = np.where(corners)
32     return x, y

```

1. **Obtaining Image Gradient** First, I obtained the image gradient for horizontal and vertical directions. To do this, I made a kernel of $[-1, 0, 1]$ and $[-1, 0, 1]^T$

For each kernel, I applied to the image with `cv2.filter2D` function. This produces I_x and I_y , each a gradient array for x direction and y direction.

2. **Calculating Cornerness and extracting corner** With Harris Corner Detector equation and the gradient, I produced the cornerness value. For hyperparameter sigma, alpha, and corner threshold was used.

$threshold \times (biggest_cornerness_value)$. If I increase the threshold, more corner points were extracted. I chose the appropriate threshold based on the given corner points count for each images.

3. **Non-Max Suppression** To apply non-max suppression, I iterated through the corner matrix, and for each pixels that are chosen for corner I checked out if it is the local maximum in 3×3 neighborhood. If it is not, I excluded it from the corner list.

SIFT Feature Descriptor

```

1  def get_descriptors(image, x, y, descriptor_window_image_width):
2
3      x = np.round(x).astype(int)
4      y = np.round(y).astype(int)
5
6      num_bins = 8
7      window_width = descriptor_window_image_width // 4 # width of
the 4x4 cells in a window
8
9      descriptors = np.zeros((len(x), 128)) # 128 = 16 histograms *
8 bins per histogram
10
11      gradient_filter_x = np.array([[ -1, 0, 1]])
12      gradient_filter_y = np.array([[ -1, 0, 1]])
13
14      # Compute gradients by convolving gradient filters with the
image
15      dx = cv2.filter2D(src=image, kernel=gradient_filter_x, ddepth
=-1)
16      dy = cv2.filter2D(src=image, kernel=gradient_filter_y, ddepth
=-1)
17      magnitude = np.sqrt(np.add(np.square(dx), np.square(dy)))
18      orientation = np.arctan2(dy, dx) * (180 / np.pi) % 360
19
20      gaussian_window = cv2.getGaussianKernel(
descriptor_window_image_width, descriptor_window_image_width / 2)
21      gaussian_window = gaussian_window * gaussian_window.T
22
23      for idx in range(len(x)):
24          xi=x[idx]
25          yi=y[idx]
26          # Ensure the window is fully within the image bounds
27          min_x = max(xi - window_width * 2, 0)
28          max_x = min(xi + window_width * 2, image.shape[1])
29          min_y = max(yi - window_width * 2, 0)

```

```

30         max_y = min(yi + window_width * 2, image.shape[0])
31
32         window_magnitude = magnitude[min_y:max_y, min_x:max_x]
33         window_orientation = orientation[min_y:max_y, min_x:max_x]
34         # Apply the Gaussian window
35         weight_window_magnitude = window_magnitude *
gaussian_window[min_y - yi + window_width * 2:max_y - yi +
window_width * 2,
36
37         min_x - xi + window_width * 2:max_x - xi + window_width * 2]
38
39         descriptor_vector = np.zeros((4, 4, num_bins))
40
41         for i in range(4):
42             for j in range(4):
43                 subregion_w_mag = weight_window_magnitude[
window_width*i:window_width*(i+1), window_width*j:window_width*(j
+1)].flatten()
44                 subregion_orientation = window_orientation[
window_width*i:window_width*(i+1), window_width*j:window_width*(j
+1)].flatten()
45
46                 hist, _ = np.histogram(subregion_orientation, bins
=num_bins, range=(0, 360), weights=subregion_w_mag)
47                 descriptor_vector[i, j, :] = hist
48
49                 # Normalize the descriptor
50                 descriptor_vector = descriptor_vector.flatten()
51                 descriptor_vector /= (np.linalg.norm(descriptor_vector) +
1e-7) # Preventing division by zero
52
53                 # Clip values to 0.2 and re-normalize
54                 descriptor_vector = np.clip(descriptor_vector, 0, 0.2)
55                 descriptor_vector /= (np.linalg.norm(descriptor_vector) +
1e-7)
56
57                 descriptors[idx, :] = descriptor_vector
58
59         return descriptors

```

1. **Calculating Image magnitude and orientation** First, for computational efficiency, I pre-calculated the magnitude and orientation of feature vectors. To do this, with the same way used in feature detection I obtained the gradient and calculated magnitude and orientation.
2. **Gaussian Window** Furthermore I calculated the gaussian kernel with the size of the descriptor window, to assign weights to the pixel values in the descriptor.
3. **Calculating descriptor** Iterating through the corner points given, I initialized the descriptor vector by applying the pre-calculated gaussian window to the local magnitude array. I also considered the boundary condition, so started from the half size of the descriptor window.

4. **Creating Histogram** For each 4×4 cell in the descriptor window, calculated the magnitude and orientation of the subregion. Then, with `np.histogram` function made a histogram for each subregions. This leads to $4 \times 4 \times 8 = 128$ dimensions feature descriptor
5. **Normalization and Clipping** Finally, I normalized the descriptor and clipped the values to 0,2; then re-normalized to obtain the final descriptor.

Feature Matching

```

1  def match_features(features1, features2):
2      matches = []
3      confidences = []
4      ratio = 0.8
5
6      for i in range(features1.shape[0]):
7
8          distances = np.sqrt(np.square(np.subtract(features1[i, :],
9              features2)).sum(axis=1))
10
11         sorted_indices = np.argsort(distances)
12         closest_neighbor_index = sorted_indices[0]
13         second_closest_neighbor_index = sorted_indices[1]
14
15         nn_distance_ratio = distances[closest_neighbor_index] /
16             distances[second_closest_neighbor_index]
17
18         if nn_distance_ratio < ratio:
19             matches.append([i, closest_neighbor_index])
20             confidences.append(1.0 - nn_distance_ratio)
21
22         matches = np.array(matches)
23         confidences = np.array(confidences)
24
25         sorted_indices = np.argsort(confidences)
26         matches = matches[sorted_indices]
27         confidences = confidences[sorted_indices]
28
29     return matches, confidences

```

1. **Calculating Nearest Neighbor** For descriptor vectors obtained from SIFT, I applied Nearest neighbor with Euclidean distance to match the good features. After calculating the Euclidean distance between the feature of two images, I sorted the vectors in order of distance. Then, I calculated the ratio between first and second nearest neighbor's distance.
2. **Applying Threshold to pick the feature point** I applied threshold of 80%, and selected the matches bigger than the threshold for good matches and appended in the array.

3. **Sorting** Finally, I calculated the confidence value by extracting the ratio from 1, and sorted all the indices with confidence.

Feature Matching Results




Figure 1: Notre Dame. In this image, with 160 submitted matches the accuracy was **89.31%** and for 100 matches in decreasing confidence the accuracy was **91%**.




Figure 2: Mount Rushmore. In this image, with 226 submitted matches the accuracy was **97.35%** and for 100 matches in decreasing confidence the accuracy was **97.00%**.



Figure 3: Episcopal Gaudi. In this image, with 5 submitted matches the accuracy was **60.00%** and for 100 matches in decreasing confidence the accuracy was **3.00%**.

We can notice that the accuracy dropped in this image. I think this is because in this picture the brick pattern in the wall makes ambiguity between the feature points; we can't easily distinguish the feature points on those repetitive patterns with nearest neighbor algorithm.

Normalized Patch Descriptor VS SIFT Descriptor

This is the normalized patch descriptor that I implemented at first.

```
1  def get_descriptors(image, x, y, descriptor_window_image_width):
2      x = np.round(x).astype(int)
3      y = np.round(y).astype(int)
4
5      patch_half_width = descriptor_window_image_width // 2
6
7      descriptors = np.zeros((len(x), descriptor_window_image_width
8      ** 2))
9
10     for idx in range(len(x)):
11         xi = x[idx]
12         yi = y[idx]
13
14         min_x = max(xi - patch_half_width, 0)
15         max_x = min(xi + patch_half_width, image.shape[1])
16         min_y = max(yi - patch_half_width, 0)
17         max_y = min(yi + patch_half_width, image.shape[0])
18
19         patch = image[min_y:max_y, min_x:max_x]
20         resized_patch = cv2.resize(patch, (
21         descriptor_window_image_width, descriptor_window_image_width))
22
23         patch_descriptor = resized_patch.flatten()
24         patch_descriptor = patch_descriptor / np.linalg.norm(
25         patch_descriptor)
26
27         descriptors[idx, :] = patch_descriptor
28
29     return descriptors
```

I simply made a patch around each feature points from Harris corner detector and normalized them and used them as the feature vector. For Notre Dame image, with this descriptor and all the other code identical, showed **78.05%** accuracy in all matches and **64%** accuracy for 100 matches by decreasing confidences, compared to **89.31%** and **91%** of SIFT descriptor.